

Data Science II - Homework 1

Bowen Xia (bx2232)

```

library(glmnet)
library(caret)
library(pls)
library(plotmo)
library(corrplot)
library(tidyverse)

train <- read.csv("housing_training.csv")
test  <- read.csv("housing_test.csv")

fac_vars <- c("Overall_Qual", "Kitchen_Qual", "Fireplace_Qu", "Exter_Qual")
train[fac_vars] <- lapply(train[fac_vars], as.factor)
test[fac_vars]  <- lapply(test[fac_vars], as.factor)

y      <- train$Sale_Price
y.test <- test$Sale_Price

x      <- model.matrix(Sale_Price ~ ., train)[, -1]
x.test <- model.matrix(Sale_Price ~ ., test)[, -1]

```

(a) Lasso


```

ctrl1 <- trainControl(method = "cv", number = 10)

set.seed(2)
lasso.fit <- train(Sale_Price ~ .,
                     data      = train,
                     method    = "glmnet",
                     tuneGrid  = expand.grid(alpha  = 1,
                                            lambda = exp(seq(10, 0, length = 100))),
                     trControl = ctrl1)

plot(lasso.fit, xTrans = log)

```



```

lasso.fit$bestTune

##      alpha    lambda
## 39      1 46.45034

coef(lasso.fit$finalModel, lasso.fit$bestTune$lambda)

## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               s=46.45034
## (Intercept)           -4.872802e+06
## Gr_Liv_Area            6.564839e+01
## First_Flr_SF            7.903386e-01
## Second_Flr_SF           .
## Total_Bsmt_SF           3.535970e+01
## Low_Qual_Fin_SF         -4.121438e+01
## Wood_Deck_SF             1.174410e+01
## Open_Porch_SF             1.565541e+01
## Bsmt_Unf_SF              -2.089048e+01
## Mas_Vnr_Area             1.077520e+01
## Garage_Cars               4.122437e+03
## Garage_Area                8.065859e+00
## Year_Built                  3.238794e+02
## TotRms_AbvGrd              -3.679193e+03
## Full_Bath                   -3.977873e+03
## Overall_QualAverage        -4.903816e+03

```

```

## Overall_QualBelow_Average -1.255837e+04
## Overall_QualExcellent 7.478188e+04
## Overall_QualFair -1.086979e+04
## Overall_QualGood 1.216852e+04
## Overall_QualVery_Excellent 1.343708e+05
## Overall_QualVery_Good 3.792662e+04
## Kitchen_QualFair -2.534577e+04
## Kitchen_QualGood -1.765462e+04
## Kitchen_QualTypical -2.572670e+04
## Fireplaces 1.077660e+04
## Fireplace_QuFair -7.718858e+03
## Fireplace_QuGood .
## Fireplace_QuNo_Fireplace 1.814710e+03
## Fireplace_QuPoor -5.691040e+03
## Fireplace_QuTypical -7.015719e+03
## Exter_QualFair -3.456536e+04
## Exter_QualGood -1.623783e+04
## Exter_QualTypical -2.066184e+04
## Lot_Frontage 1.004202e+02
## Lot_Area 6.044148e-01
## Longitude -3.344500e+04
## Latitude 5.615697e+04
## Misc_Val 8.546998e-01
## Year_Sold -5.839028e+02

# Test error
lasso.pred <- predict(lasso.fit, newdata = test)
mse.lasso <- mean((lasso.pred - y.test)^2)
mse.lasso

## [1] 441875315

# 1SE rule
lasso.res      <- lasso.fit$results
threshold     <- min(lasso.res$RMSE) +
                  lasso.res$RMSESD[which.min(lasso.res$RMSE)] / sqrt(10)
lambda.1se.caret <- max(lasso.res$lambda[lasso.res$RMSE <= threshold])

coef.1se <- coef(lasso.fit$finalModel, s = lambda.1se.caret)
n.pred   <- sum(coef.1se != 0) - 1
cat("lambda (1SE):", round(lambda.1se.caret, 4), "\n")

## lambda (1SE): 580.3529

cat("Number of predictors (1SE):", n.pred, "\n")

## Number of predictors (1SE): 35

```

The selected tuning parameter (min CV) is $\lambda = 46.4503$, with a test MSE of 4.4187531×10^8 . Under the 1SE rule ($\lambda = 580.3529$), **35 predictors** are included in the model.


(b) Elastic Net

```
set.seed(2)
enet.fit <- train(Sale_Price ~ .,
                    data      = train,
                    method    = "glmnet",
                    tuneGrid  = expand.grid(alpha = seq(0, 1, length = 21),
                                            lambda = exp(seq(10, 0, length = 100))),
                    trControl = ctrl1)

enet.fit$bestTune

##      alpha    lambda
## 357  0.15 286.1642

myCol <- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
               superpose.line   = list(col = myCol))
plot(enet.fit, par.settings = myPar, xTrans = log)
```



```
coef(enet.fit$finalModel, enet.fit$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
```

```

##                                     s=286.1642
## (Intercept)                   -4.967207e+06
## Gr_Liv_Area                  4.405619e+01
## First_Flr_SF                 2.189055e+01
## Second_Flr_SF                2.083174e+01
## Total_Bsmt_SF                3.516175e+01
## Low_Qual_Fin_SF              -2.036525e+01
## Wood_Deck_SF                 1.199785e+01
## Open_Porch_SF                1.620836e+01
## Bsmt_Unf_SF                  -2.081728e+01
## Mas_Vnr_Area                 1.120355e+01
## Garage_Cars                  4.085804e+03
## Garage_Area                  8.461457e+00
## Year_Built                    3.216183e+02
## TotRms_AbvGrd                -3.548411e+03
## Full_Bath                     -3.809514e+03
## Overall_QualAverage          -4.997161e+03
## Overall_QualBelow_Average    -1.261057e+04
## Overall_QualExcellent         7.538235e+04
## Overall_QualFair             -1.113660e+04
## Overall_QualGood              1.207073e+04
## Overall_QualVery_Excellent   1.355886e+05
## Overall_QualVery_Good        3.778627e+04
## Kitchen_QualFair             -2.456740e+04
## Kitchen_QualGood              -1.695293e+04
## Kitchen_QualTypical          -2.501140e+04
## Fireplaces                    1.078265e+04
## Fireplace_QuFair             -7.834379e+03
## Fireplace_QuGood              .
## Fireplace_QuNo_Fireplace     1.714726e+03
## Fireplace_QuPoor              -5.785251e+03
## Fireplace_QuTypical          -7.051527e+03
## Exter_QualFair                -3.344347e+04
## Exter_QualGood                -1.509552e+04
## Exter_QualTypical             -1.961356e+04
## Lot_Frontage                  1.001034e+02
## Lot_Area                      6.036008e-01
## Longitude                     -3.405942e+04
## Latitude                      5.664338e+04
## Misc_Val                       8.552979e-01
## Year_Sold                      -5.744607e+02

```

```

enet.pred <- predict(enet.fit, newdata = test)
mse.enet <- mean((enet.pred - y.test)^2)
mse.enet

```

```

## [1] 439998442

```


The selected tuning parameters are $\alpha = 0.15$ and $\lambda = 286.1642$, with a test MSE of 4.3999844×10^8 .

The 1SE rule is **not applicable** to elastic. Elastic net has two tuning parameters (α, λ).

(c) Partial Least Squares

```
set.seed(2)
pls.fit <- train(Sale_Price ~ .,
                  data      = train,
                  method    = "pls",
                  tuneGrid  = data.frame(ncomp = 1:20),
                  trControl = ctrl1,
                  preProcess = c("center", "scale"))

plot(pls.fit)
```



```
pls.fit$bestTune

##      ncomp
## 16     16

pls.pred <- predict(pls.fit, newdata = test)
mse.pls  <- mean((pls.pred - y.test)^2)
mse.pls

## [1] 446775692
```

The PLS model includes **16 components**, with a test MSE of 4.4677569×10^8 .


(d) Model Comparison

```
set.seed(2)
lm.fit <- train(Sale_Price ~ .,
                 data      = train,
                 method    = "lm",
                 trControl = ctrl1)

resamp <- resamples(list(lasso = lasso.fit,
                         enet   = enet.fit,
                         pls    = pls.fit,
                         lm     = lm.fit))
summary(resamp)

##
## Call:
## summary.resamples(object = resamp)
##
## Models: lasso, enet, pls, lm
## Number of resamples: 10
##
## MAE
##           Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## lasso 14526.50 15737.08 16601.22 16598.73 17548.67 18443.98 0
## enet   14478.62 15705.19 16576.39 16578.81 17526.31 18436.99 0
## pls    14585.07 15807.81 16639.83 16629.39 17570.18 18442.95 0
## lm     14586.71 15828.63 16641.19 16639.99 17568.06 18439.70 0
##
## RMSE
##           Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## lasso 19602.12 21578.44 23397.99 22793.99 24058.09 24945.08 0
## enet   19554.72 21551.89 23352.35 22792.47 24093.38 24961.37 0
## pls    19736.20 21657.28 23413.90 22807.42 23936.71 24902.72 0
## lm     19750.89 21655.50 23404.48 22839.27 23982.43 24990.50 0
##
## Rsquared
##           Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## lasso 0.8736637 0.8881870 0.9058714 0.9039043 0.9195235 0.9265603 0
## enet   0.8739298 0.8878946 0.9061908 0.9039321 0.9193882 0.9264909 0
## pls    0.8730979 0.8894494 0.9055091 0.9038628 0.9194429 0.9257294 0
## lm     0.8733064 0.8891481 0.9052586 0.9036505 0.9195059 0.9248318 0

bwplot(resamp, metric = "RMSE")
```



```
data.frame(
  Model      = c("Lasso", "Elastic Net", "PLS"),
  Test_MSE = round(c(mse.lasso, mse.enet, mse.pls), 2)
) |> knitr::kable()
```

Model	Test_MSE
Lasso	441875315
Elastic Net	439998442
PLS	446775692


The **elastic net** is the best model. It has the lowest test error while providing flexibility over Lasso through the additional α parameter.

(e) Lasso: caret vs glmnet

```
set.seed(2)
cv.lasso <- cv.glmnet(x, y,
                      alpha = 1,
                      lambda = exp(seq(10, 0, length = 100)))

plot(cv.lasso)
```

```
abline(h = (cv.lasso$cvm + cv.lasso$cvsd)[which.min(cv.lasso$cvm)],
       col = 4, lwd = 2)
```




```
cv.lasso$lambda.min
```

```
## [1] 51.38745
```

```
cv.lasso$lambda.1se
```

```
## [1] 785.772
```

```
plot_glmnet(cv.lasso$glmnet.fit)
```



```
predict(cv.lasso, s = "lambda.min", type = "coefficients")
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               lambda.min
## (Intercept)           -4.858436e+06
## Gr_Liv_Area          6.558518e+01
## First_Flr_SF         7.943481e-01
## Second_Flr_SF        .
## Total_Bsmt_SF        3.537392e+01
## Low_Qual_Fin_SF     -4.115127e+01
## Wood_Deck_SF         1.171129e+01
## Open_Porch_SF        1.559229e+01
## Bsmt_Unf_SF          -2.088779e+01
## Mas_Vnr_Area         1.080038e+01
## Garage_Cars           4.107153e+03
## Garage_Area           8.108573e+00
## Year_Built            3.238466e+02
## TotRms_AbvGrd        -3.662693e+03
## Full_Bath              -3.944089e+03
## Overall_QualAverage   -4.891844e+03
## Overall_QualBelow_Average -1.253195e+04
## Overall_QualExcellent  7.492544e+04
## Overall_QualFair      -1.083724e+04
## Overall_QualGood       1.214989e+04
## Overall_QualVery_Excellent 1.346747e+05
```

```

## Overall_QualVery_Good      3.790534e+04
## Kitchen_QualFair         -2.526799e+04
## Kitchen_QualGood          -1.758893e+04
## Kitchen_QualTypical       -2.567062e+04
## Fireplaces                  1.071297e+04
## Fireplace_QuFair          -7.704585e+03
## Fireplace_QuGood            .
## Fireplace_QuNo_Fireplace   1.712992e+03
## Fireplace_QuPoor           -5.678637e+03
## Fireplace_QuTypical        -7.014304e+03
## Exter_QualFair              -3.423120e+04
## Exter_QualGood               1.592280e+04
## Exter_QualTypical           -2.035264e+04
## Lot_Frontage                 1.002152e+02
## Lot_Area                      6.044007e-01
## Longitude                     -3.329199e+04
## Latitude                       5.585746e+04
## Misc_Val                      8.479681e-01
## Year_Sold                     -5.777901e+02

pred.glmnet <- as.numeric(predict(cv.lasso, newx = x.test, s = "lambda.min"))
mse.glmnet  <- mean((pred.glmnet - y.test)^2)

data.frame(
  Package  = c("caret", "glmnet"),
  Lambda   = round(c(lasso.fit$bestTune$lambda, cv.lasso$lambda.min), 4),
  Test_MSE = round(c(mse.lasso, mse.glmnet), 2)
) |> knitr::kable()

```

Package	Lambda	Test_MSE
caret	46.4503	441875315
glmnet	51.3874	441403812

`caret` searches only the user-supplied grid; While `glmnet` generates its own data-driven sequence starting from λ_{\max} . Despite these difference, both approaches yield similar test errors